


# LoCoMo / Quality and recorded latency

Qwen3.5-9B · Same 1,540 questions · 10 configurations with available timing

## BASELINES · 6 recorded configurations



Recorded mean seconds / question → faster

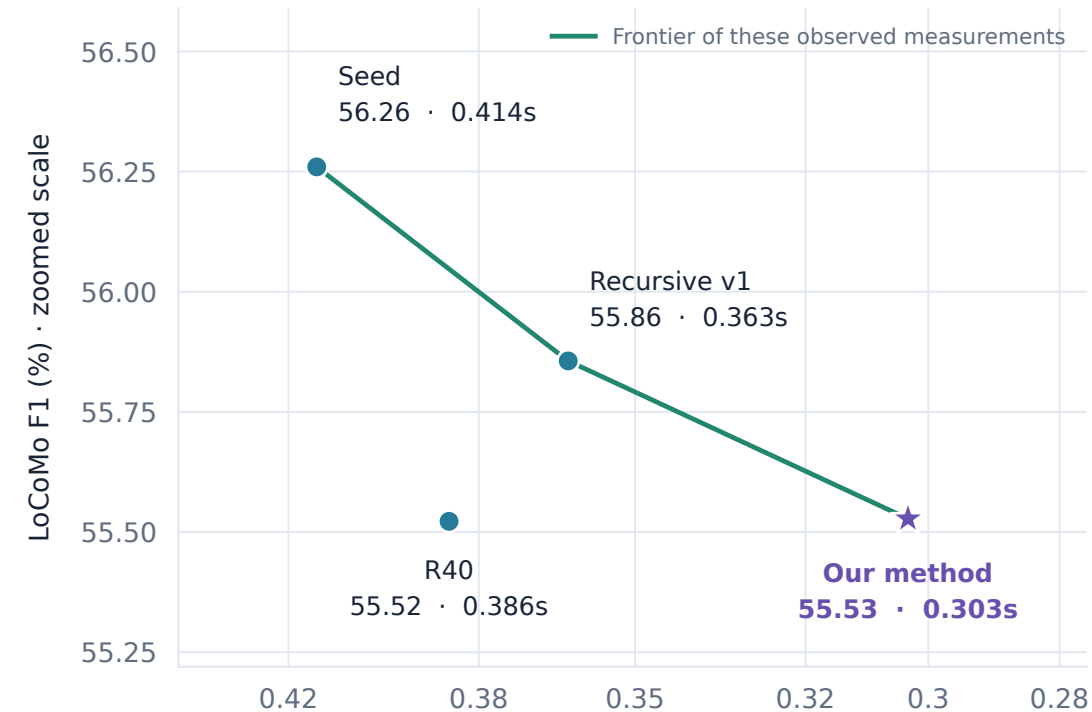
**Timing scopes differ, so the panels do not share a latency frontier.**

Baseline times follow each saved implementation; LightMem sums recorded retrieval and answer API components.

Our variants use observed LLM + embedding call time / 1,540, with batching and cache reuse; this is not end-to-end query latency.

End-to-end question times are unavailable for Our method, Full context and HiGMem. F1 ranges differ between panels.

## OUR VARIANTS · 4 recorded configurations



Amortized model-call seconds / question → faster